

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: data= pd.read_csv(r"H:\Data\.....\bank+marketing\bank\bank-full.csv")
```

Dropping unrelated columns.

```
In [3]: data= data.drop('contact', axis=1)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   day         45211 non-null  int64
9   month       45211 non-null  object
10  duration    45211 non-null  int64
11  campaign    45211 non-null  int64
12  pdays       45211 non-null  int64
13  previous    45211 non-null  int64
14  poutcome    45211 non-null  object
15  y           45211 non-null  object
dtypes: int64(7), object(9)
memory usage: 5.5+ MB
```

```
In [5]: data.head()
```

Out[5]:

	age	job	marital	education	default	balance	housing	loan	day	month	duration
0	58	management	married	tertiary	no	2143	yes	no	5	may	2
1	44	technician	single	secondary	no	29	yes	no	5	may	1
2	33	entrepreneur	married	secondary	no	2	yes	yes	5	may	
3	47	blue-collar	married	unknown	no	1506	yes	no	5	may	
4	33	unknown	single	unknown	no	1	no	no	5	may	1

```
In [6]: columns= list(data.columns)
```

```
In [7]: from sklearn.preprocessing import LabelEncoder  
le= LabelEncoder()
```

```
In [8]: for col in columns:  
        if data[col].dtype == 'object':  
            data[col]= le.fit_transform(data[col])
```

```
In [9]: data
```

Out[9]:

	age	job	marital	education	default	balance	housing	loan	day	month	duration
0	58	4	1	2	0	2143	1	0	5	8	261
1	44	9	2	1	0	29	1	0	5	8	151
2	33	2	1	1	0	2	1	1	5	8	76
3	47	1	1	3	0	1506	1	0	5	8	92
4	33	11	2	3	0	1	0	0	5	8	198
...	...	...	...	...	...	...	...	...	...	...	...
45206	51	9	1	2	0	825	0	0	17	9	977
45207	71	5	0	0	0	1729	0	0	17	9	456
45208	72	5	1	1	0	5715	0	0	17	9	1127
45209	57	1	1	1	0	668	0	0	17	9	508
45210	37	2	1	1	0	2971	0	0	17	9	361

45211 rows × 16 columns



```
In [10]: data.isna().sum()
```

```
Out[10]: age          0  
job            0  
marital        0  
education      0  
default        0  
balance        0  
housing        0  
loan           0  
day            0  
month          0  
duration       0  
campaign       0  
pdays         0  
previous       0  
poutcome       0  
y              0  
dtype: int64
```

```
In [11]: data.describe()
```

```
Out[11]:
```


	age	job	marital	education	default	balance
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.936210	4.339762	1.167725	1.224813	0.018027	1362.272058
std	10.618762	3.272657	0.608230	0.747997	0.133049	3044.765829
min	18.000000	0.000000	0.000000	0.000000	0.000000	-8019.000000
25%	33.000000	1.000000	1.000000	1.000000	0.000000	72.000000
50%	39.000000	4.000000	1.000000	1.000000	0.000000	448.000000
75%	48.000000	7.000000	2.000000	2.000000	0.000000	1428.000000
max	95.000000	11.000000	2.000000	3.000000	1.000000	102127.000000

The positive correlation (+0.39) between "duration" and "y" suggests that longer conversations during the last contact are associated with a higher likelihood of the client subscribing to the term deposit.

Note: 1. The variable "y" indicates whether the client subscribed to a term deposit (binary outcome: "yes" or "no").

2. "duration" refers to the length of the last contact with a client during a marketing campaign, measured in seconds.

```
In [12]: plt.figure(figsize=(14,7))
sns.heatmap(data.corr(), cmap= 'YlGnBu', annot= True)
plt.show()
```



```
In [13]: x= data.iloc[:, :-1]
y= data['y']
print(x.shape)
print(y.shape)
```

```
(45211, 15)
(45211,)
```

```
In [14]: from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [15]: x_train, x_test, y_train, y_test= train_test_split(x, y, train_size= 0.70, r
```

```
In [16]: dt= DecisionTreeClassifier()
dt.fit(x_train, y_train)
```

```
Out[16]: DecisionTreeClassifier()
```

```
In [17]: preds= dt.predict(x_test)
```

```
In [18]: acc= accuracy_score(preds, y_test)
print(f'Accuracy score:', acc)
```

```
Accuracy score: 0.870023591860808
```

```
In [19]: confusion_matrix(preds, y_test)
```

```
Out[19]: array([[11065,   862],
               [   901,   736]], dtype=int64)
```

```
In [20]: depths= [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15]
df= pd.DataFrame(columns= ['Max_depths', 'Accuracy'])
df
```

```
Out[20]:
```

<u>Max_depths</u>	<u>Accuracy</u>
-------------------	-----------------

```
In [21]: for input_parameter in depths:
          model= DecisionTreeClassifier(max_depth= input_parameter)
          model.fit(x_train, y_train)
          Mpreds= model.predict(x_test)
          Macc= accuracy_score(Mpreds, y_test)*100
          df= df.append({'Max_depths': input_parameter, 'Accuracy': Macc}, ignore
df
```

Out[21]:

	Max_depths	Accuracy
0	1.0	88.218815
1	2.0	88.712769
2	3.0	88.823356
3	4.0	89.184606
4	5.0	89.567974
5	6.0	89.663816
6	7.0	89.877617
7	8.0	89.752285
8	9.0	89.634326
9	10.0	89.450015
10	11.0	89.354173
11	12.0	89.332055
12	13.0	89.177234
13	14.0	88.933943
14	15.0	88.602182

The decision tree classifier achieved optimal performance on this dataset when the maximum depth was set to 7.0.